

## CODE

```
import pandas as pd

import numpy as np

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression

from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, classification_report

from sklearn.preprocessing import LabelEncoder

import seaborn as sns

import matplotlib.pyplot as plt


#load the dataset

data=pd.read_csv('titanic_sample.csv')

print(data)


#explore the dataset

print("Dataset Head:\n",data.head())

print("\nDataset Info",data.info())

print("\nMissing Values:\n",data.isnull().sum())


#Visualize some features

sns.countplot(x='Survived',data=data)

plt.title('Survival Count')

plt.show()


sns.countplot(x='Pclass',hue='Survived',data=data)

plt.title('Survival by ticket class')

plt.show()
```

```
sns.histplot(data[data['Survived']==1]['Age'],kde=True,label='Survived',color='green')
sns.histplot(data[data['Survived']==0]['Age'],kde=True,label='Not Survived',color='red')
plt.legend()
plt.title("Age Distribution by survival")
plt.show()
```

```
#data cleaning
```

```
#fill missing Age values with the median
```

```
data['Age'].fillna(data['Age'].median(),inplace=True)
```

```
print(data['Age'])
```

```
#filling missing Embarked values with the mode
```

```
data['Embarked'].fillna(data['Embarked'].mode()[0],inplace=True)
```

```
print(data['Embarked'])
```

```
#drop cabin(too many missin values)
```

```
data.drop('Cabin',axis=1,inplace=True)
```

```
print(data.drop)
```

```
#Drop irrelevant features
```

```
data.drop(['Name','Ticket','PassengerId'],axis=1,inplace=True)
```

```
print(data.drop)
```

```
#Encode categorical variables
```

```
encoder=LabelEncoder()
```

```
data['Sex']=encoder.fit_transform(data['Sex'])
```

```
data['Embarked']=encoder.fit_transform(data['Embarked'])
```

```
print(data['Sex'])
```

```
print(data['Embarked'])
```

```
#Feature Engineering
```

```
#Create familySize feature
```

```
data['FamilySize']=data['SibSp'] + data['Parch'] +1
```

```
print(data['FamilySize'])
```

```
data.drop(['SibSp','Parch'],axis=1,inplace=True)
```

```
print(data.drop)
```

```
#split the data into features and target
```

```
x=data.drop('Survived',axis=1)
```

```
print(x)
```

```
y=data['Survived']
```

```
print(y)
```

```
#train-test split
```

```
x_train, x_test, y_train, y_test=train_test_split(x,y,test_size=0.2,random_state=42)
```

```
print(x_train,x_test,y_train,y_test)
```

```
#Build and train models
```

```
#Logistic Regression
```

```
log_reg=LogisticRegression(max_iter=1000)
```

```
print(log_reg)
```

```
log_reg.fit(x_train,y_train)
```

```
print(log_reg.fit)
```

```
#Random Forest
```

```
rf=RandomForestClassifier(random_state=42)
```

```
rf.fit(x_train,y_train)
```

```
print(rf)

print(rf.fit)

from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score,
classification_report

from sklearn.ensemble import RandomForestClassifier

from sklearn.linear_model import LogisticRegression

# Assume x_train, x_test, y_train, y_test are already defined

# Define models

log_reg = LogisticRegression(random_state=42)

rf = RandomForestClassifier(random_state=42)

# Train models

log_reg.fit(x_train, y_train)

rf.fit(x_train, y_train)

# Evaluate models

models = {'Logistic Regression': log_reg, 'Random Forest': rf}

for name, model in models.items():

    y_pred = model.predict(x_test)

    print(f"\n{name} Metrics:")

    print("Accuracy:", accuracy_score(y_test, y_pred))

    print("Precision:", precision_score(y_test, y_pred))

    print("Recall:", recall_score(y_test, y_pred))

    print("F1 Score:", f1_score(y_test, y_pred))

    print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

```
# Optimization: Tune Random Forest

tuned_rf = RandomForestClassifier(n_estimators=100, max_depth=5, random_state=42)

tuned_rf.fit(x_train, y_train)

y_pred_tuned = tuned_rf.predict(x_test) # Predict on x_test, not x_train


print("\nTuned Random Forest Metrics:")

print("Accuracy:", accuracy_score(y_test, y_pred_tuned))

print("Precision:", precision_score(y_test, y_pred_tuned))

print("Recall:", recall_score(y_test, y_pred_tuned))

print("F1 Score:", f1_score(y_test, y_pred_tuned))
```

## OUTPUT

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	
0	1	0	3	Braund, Mr. Owen Harris	male	22	1	0	A/5 21171
7.2500	NaN	S							
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38	1	0	PC
17599	71.2833	C85	C						
2	3	1	3	Heikkinen, Miss. Laina	female	26	0	0	STON/O2. 3101282
7.9250	NaN	S							
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35	1	0	113803
53.1000	C123	S							
4	5	0	3	Allen, Mr. William Henry	male	35	0	0	373450
8.0500	NaN	S							

Dataset Head:

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	
0	1	0	3	Braund, Mr. Owen Harris	male	22	1	0	A/5 21171
7.2500	NaN	S							
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38	1	0	PC
17599	71.2833	C85	C						

2	3	1	3	Heikkinen, Miss. Laina	female	26	0	0	STON/O2. 3101282
7.9250	NaN	S							
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35	1	0	113803
53.1000	C123	S							
4	5	0	3	Allen, Mr. William Henry	male	35	0	0	373450 8.0500
NaN	S								

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 5 entries, 0 to 4

Data columns (total 12 columns):

#	Column	Non-Null Count	Dtype
---	--------	----------------	-------

--- -----

0	PassengerId	5 non-null	int64
1	Survived	5 non-null	int64
2	Pclass	5 non-null	int64
3	Name	5 non-null	object
4	Sex	5 non-null	object
5	Age	5 non-null	int64
6	SibSp	5 non-null	int64
7	Parch	5 non-null	int64
8	Ticket	5 non-null	object
9	Fare	5 non-null	float64
10	Cabin	2 non-null	object
11	Embarked	5 non-null	object

dtypes: float64(1), int64(6), object(5)


memory usage: 612.0+ bytes

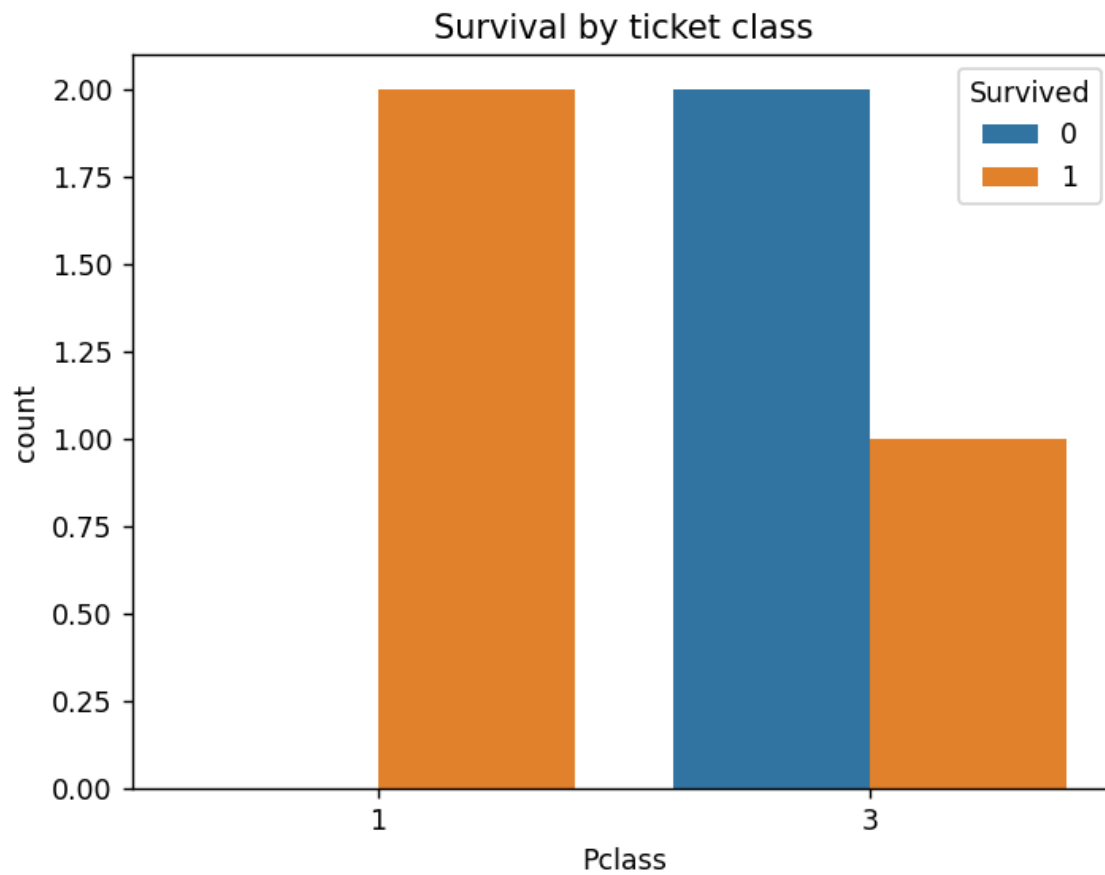
Dataset Info None

Missing Values:

PassengerId 0

Survived 0  
Pclass 0  
Name 0  
Sex 0  
Age 0  
SibSp 0  
Parch 0  
Ticket 0  
Fare 0  
Cabin 3  
Embarked 0  
dtype: int64





```
data['Age'].fillna(data['Age'].median(),inplace=True)
```

```
0 22
```

```
1 38
```

```
2 26
```

```
3 35
```

```
4 35
```

```
Name: Age, dtype: int64
```

```
data['Embarked'].fillna(data['Embarked'].mode()[0],inplace=True)
```

```
0 S
```



1 C

2 S

3 S

4 S

Name: Embarked, dtype: object

<bound method DataFrame.drop of PassengerId Survived Pclass

	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Embarked
0		1	0	3		Braund, Mr. Owen Harris	male 22 1 0 A/5 21171 7.2500	S
1		2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female 38 1 0 PC 17599 71.2833	C	
2		3	1	3	Heikkinen, Miss. Laina	female 26 0 0 STON/O2. 3101282 7.9250	S	
3		4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female 35 1 0 113803 53.1000	S	
4		5	0	3	Allen, Mr. William Henry	male 35 0 0 373450 8.0500 S>		

<bound method DataFrame.drop of Survived Pclass Sex Age SibSp Parch Fare Embarked

0	0	3	male	22	1	0	7.2500	S
1	1	1	female	38	1	0	71.2833	C
2	1	3	female	26	0	0	7.9250	S
3	1	1	female	35	1	0	53.1000	S
4	0	3	male	35	0	0	8.0500	S>

0 1

1 0

2 0

3 0

4 1

Name: Sex, dtype: int64

0 1

1 0

2 1

3 1

4 1

Name: Embarked, dtype: int64

0 2

1 2

2 1

3 2

4 1

Name: FamilySize, dtype: int64

<bound method DataFrame.drop of Survived Pclass Sex Age Fare Embarked FamilySize

0 0 3 1 22 7.2500 1 2

1 1 1 0 38 71.2833 0 2

2 1 3 0 26 7.9250 1 1

3 1 1 0 35 53.1000 1 2

4 0 3 1 35 8.0500 1 1>

Pclass Sex Age Fare Embarked FamilySize

0 3 1 22 7.2500 1 2

1 1 0 38 71.2833 0 2

2 3 0 26 7.9250 1 1

3 1 0 35 53.1000 1 2

4 3 1 35 8.0500 1 1

0 0

1 1

2 1

3 1

4 0

Name: Survived, dtype: int64

	Pclass	Sex	Age	Fare	Embarked	FamilySize	
4	3	1	35	8.050	1	1	
2	3	0	26	7.925	1	1	
0	3	1	22	7.250	1	2	
3	1	0	35	53.100	1	2	Pclass Sex Age Fare Embarked FamilySize
1	1	0	38	71.2833	0	2	4 0
2	1						
0	0						
3	1						

Name: Survived, dtype: int64 1 1

Name: Survived, dtype: int64

LogisticRegression(max\_iter=1000)

<bound method LogisticRegression.fit of LogisticRegression(max\_iter=1000)>

RandomForestClassifier(random\_state=42)

<bound method BaseForest.fit of RandomForestClassifier(random\_state=42)>

Logistic Regression Metrics:

Accuracy: 1.0

Precision: 1.0

Recall: 1.0

F1 Score: 1.0

Classification Report:

	precision	recall	f1-score	support
1	1.00	1.00	1.00	1
accuracy			1.00	1
macro avg	1.00	1.00	1.00	1

weighted avg	1.00	1.00	1.00	1
--------------	------	------	------	---

#### Random Forest Metrics:

Accuracy: 1.0

Precision: 1.0

Recall: 1.0

F1 Score: 1.0

#### Classification Report:

	precision	recall	f1-score	support
1	1.00	1.00	1.00	1
accuracy			1.00	1
macro avg	1.00	1.00	1.00	1
weighted avg	1.00	1.00	1.00	1

#### Tuned Random Forest Metrics:

Accuracy: 1.0

Precision: 1.0

Recall: 1.0

F1 Score: 1.0

